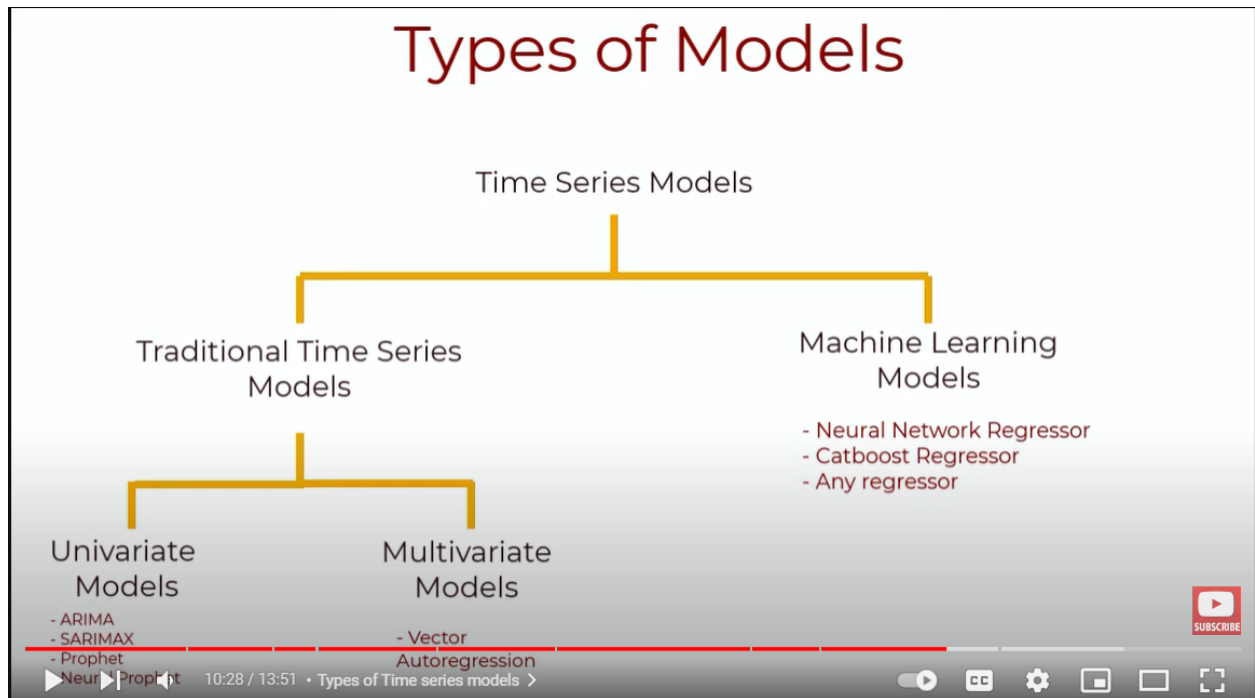


# Time series

Traditional Time Series Techniques:




## ② One Sample t-test

Z-test  $\Rightarrow$  population std

t-test  $\Rightarrow$  unknown population std

# Time series



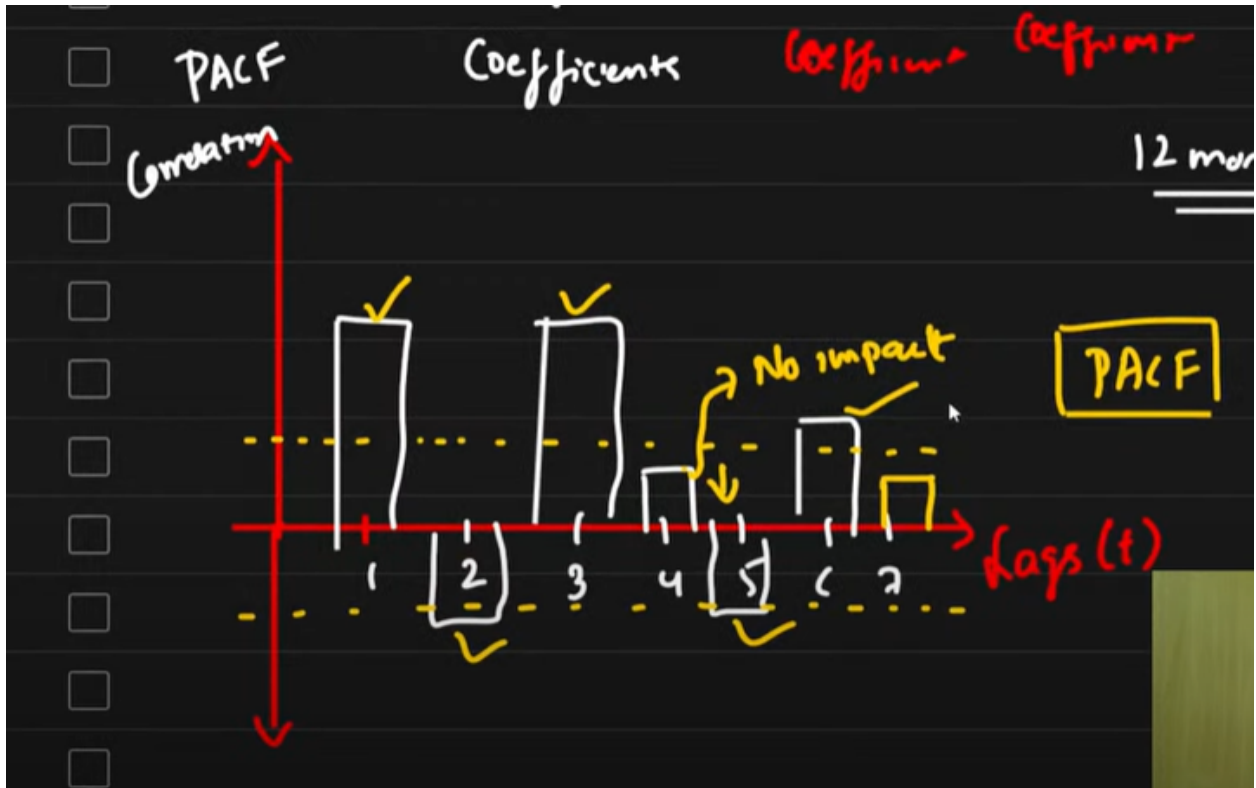
Auto Regression made by **PACF** plot

Moving Average made by **ACF** plot

Integrated made by Differencing

# Time series

### PACF PLOT :



This plot helps us to find out which graph have impact on our sells , for example 4&7 have no impact but 1,2,3,5,6 have impact

**PACF in general helps us to check the correlation**

## ARIMA steps:

## ARIMA or SARIMAX {Stationary Data}

① Check the data for stationery

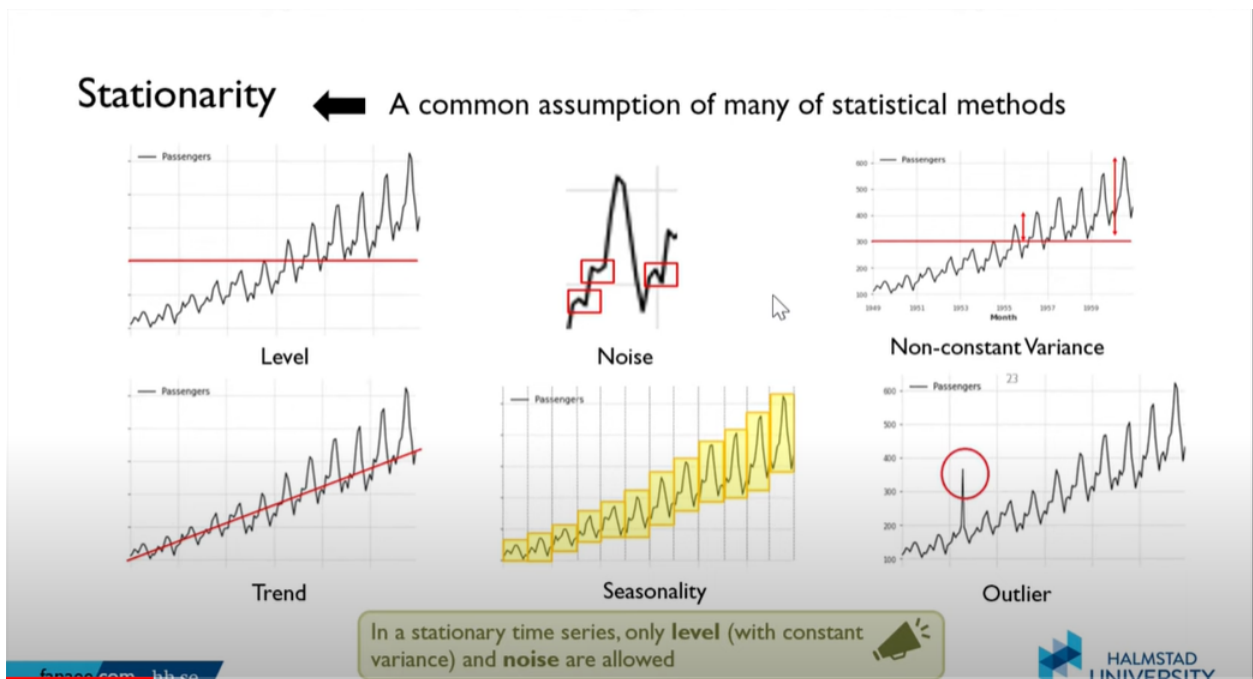
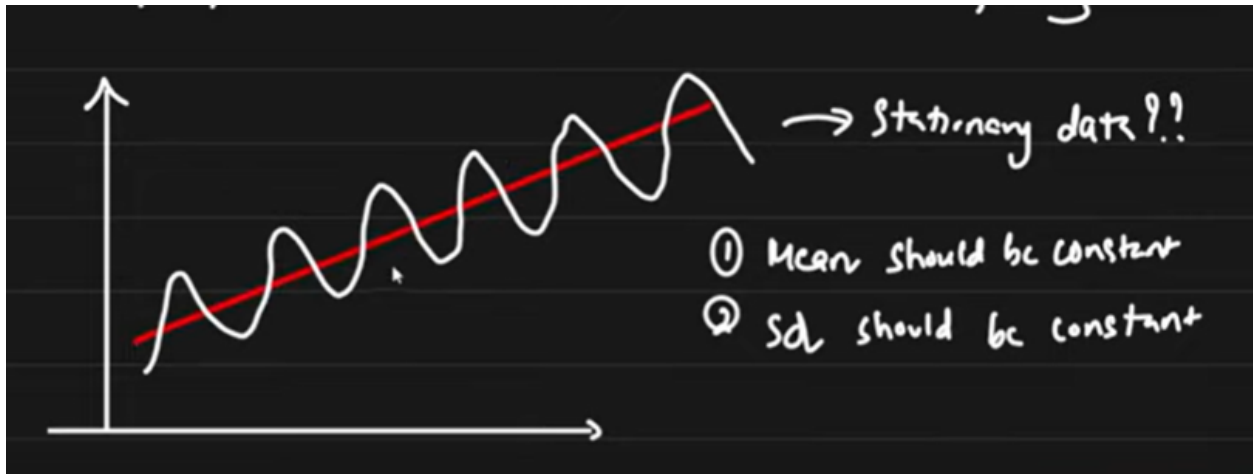
↓

Hypothesis     Dickey Fuller Test

$H_0 \Rightarrow$ Unit Root = 1 $\Rightarrow$ Data is not stationery	}
$H_1 \Rightarrow$ Unit Root < 1 $\Rightarrow$ Data is stationery	

# Time series

Check if data is stationary:



Stationary we are only allowed to have constant variance and noises


# Time series

## 1 / Stationary check :

### Stationarity Check: Statistical Tests

#### Augmented Dickey-Fuller Test

```
from pandas import read_csv
from statsmodels.tsa.stattools import adfuller
df = read_csv('airline-passengers.csv', header=0, index_col=0)
result = adfuller(df.values)
print('p-value: %f' % result[1])
```



P-value: 0.991880

P-value > 0.05: Fail to reject the null hypothesis (H0)

Null Hypothesis (H0): It looks non-stationary

Alternate Hypothesis (H1): It looks stationary

#### Implementation python :

```
def adf_test(series):
    result=adfuller(series)
    print('ADF Statistics: {}'.format(result[0]))
    print('p- value: {}'.format(result[1]))
    if result[1] <= 0.05:
        print("strong evidence against the null hypothesis, reject the null hypothesis. Data has no unit root and is stationary")
    else:
        print("weak evidence against null hypothesis, time series has a unit root, indicating it is non-stationary ")
```

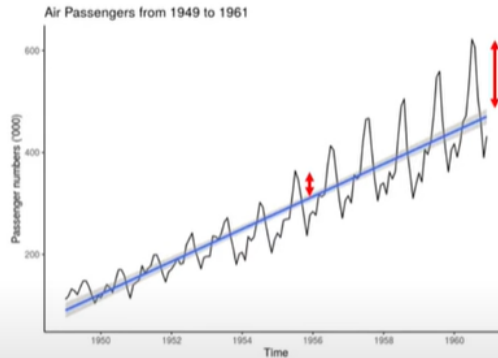
```
adf_test(df['Passengers Second Difference'].dropna())
```

# Time series

## 2 / Fixing non-constant Variance :

### Fixing Non-constant Variance

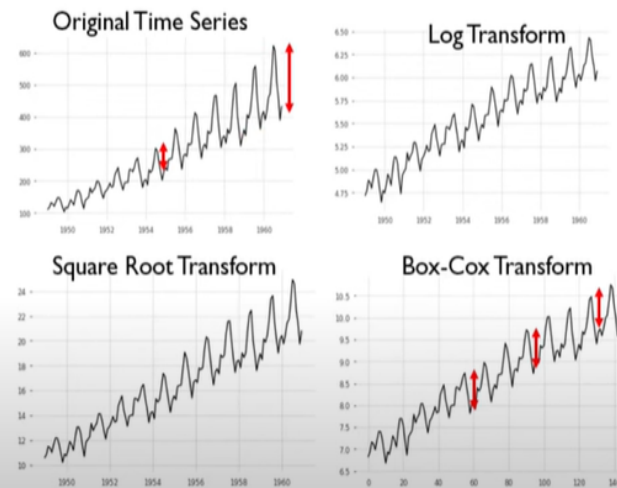
- **Logarithmic**
  - Removes extreme (exponential) trends
- **Square root transformation**
  - Removes a quadratic growth trend
- **Box-Cox transform**
  - Supports both square root and log transform



### 2-1/ Implementation

### Fixing Non-constant Variance - Python

```
from pandas import read_csv
from numpy import log
from numpy import sqrt
from scipy.stats import boxcox
import matplotlib.pyplot as plt
series = read_csv('airline-passengers.csv', header=0,
                  parse_dates=[0], index_col=0, squeeze=True)
series_log = log(series)
series_sqr = sqrt(series)
series_boxcox, lam = boxcox(series)
print('Lambda: %f % lam)
plt.plot(series_log)
plt.plot(series_sqr)
plt.plot(series_boxcox)
plt.show()
```



# Time series

## 3 / Remove trends :

### Trend Removal

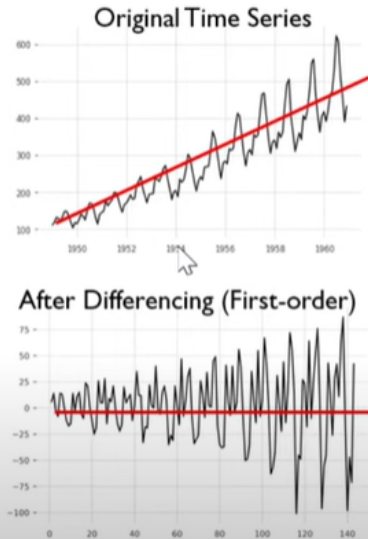
- **Differencing**

- First order : the time series value  $y_i$  is replaced by the difference between it and the previous value.

- First-order  $y'_i = y_i - y_{i-1}$ .

- Second-order  $y''_i = y'_i - y'_{i-1}$

```
from pandas import datetime
from matplotlib import pyplot
series = read_csv('airline-passengers.csv', header=0,
parse_dates=[0], index_col=0, squeeze=True)
diff = series.diff()
pyplot.plot(series.values)
pyplot.plot(diff.values)
```



### Implementation:

```
## Use Techniques Differencing
df['Passengers First Difference']=df['Thousands of
Passengers']-df['Thousands of Passengers'].shift(1)
```

This is our first shift so we can use multiple shift to remove trends so it called differencing as I in ARIMA model


# Time series

We can also combine between fixing non-constant and trend removal


## Fixing Non-constant Variance + Trend Removal

```
from pandas import read_csv
from pandas import Series
import pandas as pd
from scipy.stats import boxcox
series = read_csv('airline-passengers.csv',
header=0, parse_dates=[0], index_col=0,
squeeze=True)
series_boxcox, lam = boxcox(series)
series_boxcox = pd.DataFrame(series_boxcox)
series_boxcox_diff = series_boxcox.diff()
plt.plot(series)
plt.plot(series_boxcox_diff)
```

Original Time Series



Box Cox Transform + Differencing





## Time series

# Time Series Forecasting

### ▪ Statistical Models

#### 1 / Exponential smoothing :

Exponential smoothing is a broadly accurate forecasting method for **short-term forecasts**. The technique assigns larger weights to more recent observations while assigning exponentially decreasing weights as the observations get increasingly distant. This method produces slightly unreliable long-term forecasts.

- Simple Exponential Smoothing (Brown, 1956)
  - Prediction is a weighted linear sum of recent past observations or lags.
  - Weights decline exponentially, keeping little record of past
  - General Formulation:  $Y_t = W \cdot X_t + (1 - W) \cdot Y_{t-1}$ 
    - $W$ : Smoothing Factor

Example: [4, 6, 5, 3, 7],  $W=0.2$

t	Xt	Yt	Yt+1
t=1	X1=4	Y1=4.0	N/A
t=2	X2=6	Y2=(.2)(6) + (1-.2)(4.0) = 4.4	4.0
t=3	X3=5	Y3=(.2)(5) + (1-.2)(4.4) = 4.5	4.4
t=4	X4=3	Y4=(.2)(3) + (1-.2)(4.5) = 4.2	4.5
t=5	X5=7	Y5=(.2)(7) + (1-.2)(4.2) = 4.8	4.2

Output= [N/A, 4.0 4.4 4.5 4.2]  
Forecast for t=6 → 4.2

The latest lags have more weight to the old lags so we give more importance to the new observations. As mentioned in the general formulation.

We have also other type of smoothing such as :

- Double Exponential Smoothing (Holt, 1957)
  - Smoothing for Trend + Level
- Triple Exponential Smoothing (Holt-Winter, 1960)
  - Smoothing for Trend + Seasonality + Level

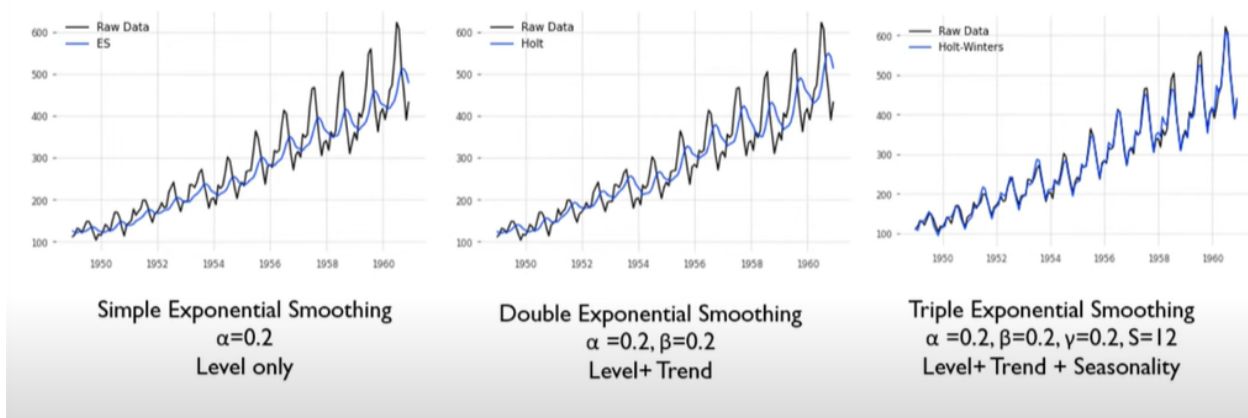
# Time series

## Exponential Smoothing - Python

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from statsmodels.tsa.holtwinters import ExponentialSmoothing
from statsmodels.tsa.holtwinters import SimpleExpSmoothing
df = pd.read_csv('airline-passengers.csv', parse_dates=['Month'], index_col='Month')
df.index.freq = 'MS'
ts=df.iloc[:, 0]
es1 = SimpleExpSmoothing(ts).fit(smoothing_level=0.2)
ts_es1 = es1.predict(start=ts.index[0], end=ts.index[-1])
es2 = ExponentialSmoothing(ts, trend='add').fit(smoothing_level=0.2, smoothing_trend=0.2)
ts_es2 = es2.predict(start=ts.index[0], end=ts.index[-1])
es3 = ExponentialSmoothing(ts, trend='add', seasonal='mul',
seasonal_periods=12).fit(smoothing_level=0.2, smoothing_trend=0.2, smoothing_seasonal=0.2)
ts_es3 = es3.predict(start=ts.index[0], end=ts.index[-1])
```

Result :

## Exponential Smoothing - Python



Here we see that triple exponential smoothing has better and relevant results than simple and double smoothing which is normal.

# Time series

## Assumption of Models

Model	Constant Variance	Constant Mean (No Trend)	Seasonality-Free
AR	✓	✓	✓
MA	✓	✓	✓
ARMA	✓	✓	✓
ARIMA	✓	X	✓
SARIMA	✓	X	X

Model selection (Hyperparameter Configuration)

## Model Selection (Hyperparameter Configuration)

- Method 1: Visualization
  - MA model: ACF plot
  - AR Model: PACF plot
- Method 2: Hyperparameter Tuning with Grid Search
  - Test for different model orders exhaustively
    - A model that is more accurate and less complex is preferred
      - More accurate: when fitting train data, it has less error
      - Less complex: has a smaller number of coefficients

# Time series

## Autocorrelation Plot (ACF)

- ACF describes the autocorrelation between an observation and another observation at a prior time step that includes direct and indirect dependence information.
- How to calculate ACF?

For lag=1:MaxLag:

$$\text{ACF}(\text{lag}) = \text{corr}(X_t, X_{\text{lag}N})$$

- Example:  $X = [1 \ 2 \ 3 \ 4 \ 5 \ 6 \ 7 \ 8]$


$$X = [3 \ 4 \ 5 \ 6 \ 7 \ 8]$$

$$X_{\text{lag}1} = [2 \ 3 \ 4 \ 5 \ 6 \ 7]$$

$$X_{\text{lag}2} = [1 \ 2 \ 3 \ 4 \ 5 \ 6]$$

$$\text{ACF}(\text{lag}=1) = \text{correlation}(X, X_{\text{lag}1})$$

$$\text{ACF}(\text{lag}=2) = \text{correlation}(X, X_{\text{lag}2})$$



Only 4 of these lags have an impact on our data.

## Advantages and Disadvantages of Statistical Models

### ▪ Advantage:

- Simple, Transparent, Rigorous
- State-of-the-art performance for small data sets better than ML/deep learning

### ▪ Disadvantages

- No more performance improvement when given larger data.
  - Not very good for very large datasets
- Unable to capture nonlinear dynamics
- One-step ahead forecast

# Time series


ANOMALY DETECTION :

1 / STL

## Seasonal-Trend Decomposition using LOESS (STL)

The basic idea is that if you have a time series with a regular pattern to it, you can run the series through the STL algorithm and isolate the regular pattern. Everything left over is by definition “irregular”, and anomaly detection amounts to deciding whether the irregularity is sufficiently large.

## Time series



The algorithm decomposes the series into three components: seasonal, trend and remainder. The *seasonal* is the periodic component, the *trend* is general up/down behavior, and the *remainder* is what's left over. The seasonal and trend taken together form the “regular” part of the series, and thus the part that we want to discount during anomaly detection.

# Time series

## **What about multiple seasonalities?**

Some time series have more than one seasonality. For example, at Expedia (where I work), bookings time series have three seasonalities: daily, weekly and annual.

While there are procedures that generate decompositions with multiple seasonal components, STL doesn't do that. The highest frequency seasonality gets to be the seasonal component, and any lower frequency seasonalities get absorbed into the trend.

2 / Isolation Forest